

Project-01-Identifying-predictors

July 3, 2022

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
%matplotlib inline
```

Load the dataset

```
[3]: df = pd.read_csv(r'C:\Users\krzys\Desktop\automobileEDA.csv')
```

Check the dataset

```
[4]: df.dtypes
```

```
[4]: symboling          int64
normalized-losses     int64
make                  object
aspiration            object
num-of-doors          object
body-style            object
drive-wheels          object
engine-location       object
wheel-base           float64
length               float64
width                float64
height               float64
curb-weight           int64
engine-type          object
num-of-cylinders      object
engine-size           int64
fuel-system          object
bore                 float64
stroke               float64
compression-ratio     float64
horsepower            float64
peak-rpm             float64
city-mpg              int64
```

```
highway-mpg          int64
price                float64
city-L/100km         float64
horsepower-binned    object
diesel               int64
gas                  int64
dtype: object
```

```
[5]: #Missing values
missing_data = df.isnull()

for column in missing_data.columns.values.tolist():
    print(column)
    print (missing_data[column].value_counts())
    print("")

#Fix missing values in stroke column
df.interpolate(inplace=True) #fix missing in stroke
df.drop(['horsepower-binned'], inplace=True, axis=1)
```

```
symboling
False    201
Name: symboling, dtype: int64
```

```
normalized-losses
False    201
Name: normalized-losses, dtype: int64
```

```
make
False    201
Name: make, dtype: int64
```

```
aspiration
False    201
Name: aspiration, dtype: int64
```

```
num-of-doors
False    201
Name: num-of-doors, dtype: int64
```

```
body-style
False    201
Name: body-style, dtype: int64
```

```
drive-wheels
False    201
Name: drive-wheels, dtype: int64
```

```

engine-location
False      201
Name: engine-location, dtype: int64

wheel-base
False      201
Name: wheel-base, dtype: int64

length
False      201
Name: length, dtype: int64

width
False      201
Name: width, dtype: int64

height
False      201
Name: height, dtype: int64

curb-weight
False      201
Name: curb-weight, dtype: int64

engine-type
False      201
Name: engine-type, dtype: int64

num-of-cylinders
False      201
Name: num-of-cylinders, dtype: int64

engine-size
False      201
Name: engine-size, dtype: int64

fuel-system
False      201
Name: fuel-system, dtype: int64

bore
False      201
Name: bore, dtype: int64

stroke
False      197
True        4

```

```

Name: stroke, dtype: int64

compression-ratio
False      201
Name: compression-ratio, dtype: int64

horsepower
False      201
Name: horsepower, dtype: int64

peak-rpm
False      201
Name: peak-rpm, dtype: int64

city-mpg
False      201
Name: city-mpg, dtype: int64

highway-mpg
False      201
Name: highway-mpg, dtype: int64

price
False      201
Name: price, dtype: int64

city-L/100km
False      201
Name: city-L/100km, dtype: int64

horsepower-binned
False      200
True         1
Name: horsepower-binned, dtype: int64

diesel
False      201
Name: diesel, dtype: int64

gas
False      201
Name: gas, dtype: int64

```

Correlations

```

[6]: #Quick look
df_corr = df.corr()

```

```
df_corr.sort_values(by=['price'])
```

```
[6]:
```

| | symboling | normalized-losses | wheel-base | length | \ |
|-------------------|-----------|-------------------|------------|-----------|---|
| highway-mpg | 0.036233 | -0.181877 | -0.543304 | -0.698142 | |
| city-mpg | -0.035527 | -0.225016 | -0.470606 | -0.665192 | |
| gas | 0.196735 | 0.101546 | -0.307237 | -0.211187 | |
| peak-rpm | 0.279740 | 0.239543 | -0.360305 | -0.285970 | |
| symboling | 1.000000 | 0.466264 | -0.535987 | -0.365404 | |
| compression-ratio | -0.182196 | -0.114713 | 0.250313 | 0.159733 | |
| stroke | -0.006564 | 0.055836 | 0.157438 | 0.123525 | |
| diesel | -0.196735 | -0.101546 | 0.307237 | 0.211187 | |
| normalized-losses | 0.466264 | 1.000000 | -0.056661 | 0.019424 | |
| height | -0.550160 | -0.373737 | 0.590742 | 0.492063 | |
| bore | -0.140019 | -0.029862 | 0.493244 | 0.608971 | |
| wheel-base | -0.535987 | -0.056661 | 1.000000 | 0.876024 | |
| length | -0.365404 | 0.019424 | 0.876024 | 1.000000 | |
| width | -0.242423 | 0.086802 | 0.814507 | 0.857170 | |
| city-L/100km | 0.066171 | 0.238567 | 0.476153 | 0.657373 | |
| horsepower | 0.075819 | 0.217299 | 0.371147 | 0.579821 | |
| curb-weight | -0.233118 | 0.099404 | 0.782097 | 0.880665 | |
| engine-size | -0.110581 | 0.112360 | 0.572027 | 0.685025 | |
| price | -0.082391 | 0.133999 | 0.584642 | 0.690628 | |

| | width | height | curb-weight | engine-size | bore | \ |
|-------------------|-----------|-----------|-------------|-------------|-----------|---|
| highway-mpg | -0.680635 | -0.104812 | -0.794889 | -0.679571 | -0.591309 | |
| city-mpg | -0.633531 | -0.049800 | -0.749543 | -0.650546 | -0.582027 | |
| gas | -0.244356 | -0.281578 | -0.221046 | -0.070779 | -0.054458 | |
| peak-rpm | -0.245800 | -0.309974 | -0.279361 | -0.256733 | -0.267392 | |
| symboling | -0.242423 | -0.550160 | -0.233118 | -0.110581 | -0.140019 | |
| compression-ratio | 0.189867 | 0.259737 | 0.156433 | 0.028889 | 0.001263 | |
| stroke | 0.188681 | -0.062214 | 0.167397 | 0.204933 | -0.055376 | |
| diesel | 0.244356 | 0.281578 | 0.221046 | 0.070779 | 0.054458 | |
| normalized-losses | 0.086802 | -0.373737 | 0.099404 | 0.112360 | -0.029862 | |
| height | 0.306002 | 1.000000 | 0.307581 | 0.074694 | 0.180449 | |
| bore | 0.544885 | 0.180449 | 0.644060 | 0.572609 | 1.000000 | |
| wheel-base | 0.814507 | 0.590742 | 0.782097 | 0.572027 | 0.493244 | |
| length | 0.857170 | 0.492063 | 0.880665 | 0.685025 | 0.608971 | |
| width | 1.000000 | 0.306002 | 0.866201 | 0.729436 | 0.544885 | |
| city-L/100km | 0.673363 | 0.003811 | 0.785353 | 0.745059 | 0.554610 | |
| horsepower | 0.615077 | -0.087027 | 0.757976 | 0.822676 | 0.566936 | |
| curb-weight | 0.866201 | 0.307581 | 1.000000 | 0.849072 | 0.644060 | |
| engine-size | 0.729436 | 0.074694 | 0.849072 | 1.000000 | 0.572609 | |
| price | 0.751265 | 0.135486 | 0.834415 | 0.872335 | 0.543155 | |

| | stroke | compression-ratio | horsepower | peak-rpm | \ |
|-------------|-----------|-------------------|------------|-----------|---|
| highway-mpg | -0.035665 | 0.268465 | -0.804575 | -0.058598 | |
| city-mpg | -0.035333 | 0.331425 | -0.822214 | -0.115413 | |

| | | | | |
|-------------------|-----------|-----------|-----------|-----------|
| gas | -0.240684 | -0.985231 | 0.169053 | 0.475812 |
| peak-rpm | -0.061840 | -0.435780 | 0.107885 | 1.000000 |
| symboling | -0.006564 | -0.182196 | 0.075819 | 0.279740 |
| compression-ratio | 0.187638 | 1.000000 | -0.214514 | -0.435780 |
| stroke | 1.000000 | 0.187638 | 0.099424 | -0.061840 |
| diesel | 0.240684 | 0.985231 | -0.169053 | -0.475812 |
| normalized-losses | 0.055836 | -0.114713 | 0.217299 | 0.239543 |
| height | -0.062214 | 0.259737 | -0.087027 | -0.309974 |
| bore | -0.055376 | 0.001263 | 0.566936 | -0.267392 |
| wheel-base | 0.157438 | 0.250313 | 0.371147 | -0.360305 |
| length | 0.123525 | 0.159733 | 0.579821 | -0.285970 |
| width | 0.188681 | 0.189867 | 0.615077 | -0.245800 |
| city-L/100km | 0.038001 | -0.299372 | 0.889488 | 0.115830 |
| horsepower | 0.099424 | -0.214514 | 1.000000 | 0.107885 |
| curb-weight | 0.167397 | 0.156433 | 0.757976 | -0.279361 |
| engine-size | 0.204933 | 0.028889 | 0.822676 | -0.256733 |
| price | 0.082982 | 0.071107 | 0.809575 | -0.101616 |

| | city-mpg | highway-mpg | price | city-L/100km | diesel \ |
|-------------------|-----------|-------------|-----------|--------------|-----------|
| highway-mpg | 0.972044 | 1.000000 | -0.704692 | -0.930028 | 0.198690 |
| city-mpg | 1.000000 | 0.972044 | -0.686571 | -0.949713 | 0.265676 |
| gas | -0.265676 | -0.198690 | -0.110326 | 0.241282 | -1.000000 |
| peak-rpm | -0.115413 | -0.058598 | -0.101616 | 0.115830 | -0.475812 |
| symboling | -0.035527 | 0.036233 | -0.082391 | 0.066171 | -0.196735 |
| compression-ratio | 0.331425 | 0.268465 | 0.071107 | -0.299372 | 0.985231 |
| stroke | -0.035333 | -0.035665 | 0.082982 | 0.038001 | 0.240684 |
| diesel | 0.265676 | 0.198690 | 0.110326 | -0.241282 | 1.000000 |
| normalized-losses | -0.225016 | -0.181877 | 0.133999 | 0.238567 | -0.101546 |
| height | -0.049800 | -0.104812 | 0.135486 | 0.003811 | 0.281578 |
| bore | -0.582027 | -0.591309 | 0.543155 | 0.554610 | 0.054458 |
| wheel-base | -0.470606 | -0.543304 | 0.584642 | 0.476153 | 0.307237 |
| length | -0.665192 | -0.698142 | 0.690628 | 0.657373 | 0.211187 |
| width | -0.633531 | -0.680635 | 0.751265 | 0.673363 | 0.244356 |
| city-L/100km | -0.949713 | -0.930028 | 0.789898 | 1.000000 | -0.241282 |
| horsepower | -0.822214 | -0.804575 | 0.809575 | 0.889488 | -0.169053 |
| curb-weight | -0.749543 | -0.794889 | 0.834415 | 0.785353 | 0.221046 |
| engine-size | -0.650546 | -0.679571 | 0.872335 | 0.745059 | 0.070779 |
| price | -0.686571 | -0.704692 | 1.000000 | 0.789898 | 0.110326 |

| | |
|-------------------|-----------|
| | gas |
| highway-mpg | -0.198690 |
| city-mpg | -0.265676 |
| gas | 1.000000 |
| peak-rpm | 0.475812 |
| symboling | 0.196735 |
| compression-ratio | -0.985231 |
| stroke | -0.240684 |

| | |
|-------------------|-----------|
| diesel | -1.000000 |
| normalized-losses | 0.101546 |
| height | -0.281578 |
| bore | -0.054458 |
| wheel-base | -0.307237 |
| length | -0.211187 |
| width | -0.244356 |
| city-L/100km | 0.241282 |
| horsepower | 0.169053 |
| curb-weight | -0.221046 |
| engine-size | -0.070779 |
| price | -0.110326 |

```
[7]: #Only the following numerical variables should be considered in further
      ↪analysis (due to moderate or strong
      #positive/negative correlation): length, width, curb-weight, engine-size,
      ↪horsepower, city-mpg, highway-mpg,
      #wheel-base, bore.
```

```
[8]: #Detailed Pearson correlation analysis


print('Length')
pearson_coef, p_value = stats.pearsonr(df['length'], df['price'])
print("Correlation Coefficient:", pearson_coef, " P-value:", p_value)
print(" ")
#Since the p-value is < 0.001, the correlation between length and price is
↪statistically
#significant, and the linear relationship is moderately strong (~0.691).

sns.regplot(x="length", y="price", data=df)
plt.ylim(0,)
```

Length

Correlation Coefficient: 0.690628380448364 P-value: 8.016477466158986e-30

```
[8]: (0.0, 47867.09963559985)
```




```
[9]: print('Width')
      pearson_coef, p_value = stats.pearsonr(df['width'], df['price'])
      print("Correlation Coefficient:", pearson_coef, " P-value:", p_value)
      print(" ")
      #Since the p-value is < 0.001, the correlation between width and price is
      ↪statistically
      #significant, and the linear relationship is quite strong (~0.751).

      sns.regplot(x="width", y="price", data=df)
      plt.ylim(0,)
```

Width

Correlation Coefficient: 0.7512653440522674 P-value: 9.200335510481516e-38

```
[9]: (0.0, 47914.503239972415)
```


```
[10]: print('Curb-weight')
pearson_coef, p_value = stats.pearsonr(df['curb-weight'], df['price'])
print( "Correlation Coefficient:", pearson_coef, " P-value:", p_value)
print(" ")
#Since the p-value is < 0.001, the correlation between curb-weight and price is
↳statistically
#significant, and the linear relationship is quite strong (~0.834).

sns.regplot(x="curb-weight", y="price", data=df)
plt.ylim(0,)
```

Curb-weight

Correlation Coefficient: 0.8344145257702846 P-value: 2.1895772388936914e-53

```
[10]: (0.0, 47754.78706739099)
```




```
[11]: print('Engine-size')
pearson_coef, p_value = stats.pearsonr(df['engine-size'], df['price'])
print("Correlation Coefficient:", pearson_coef, " P-value:", p_value)
print(" ")
#Since the p-value is < 0.001, the correlation between engine-size and price is
↳statistically
#significant, and the linear relationship is very strong (~0.872).

sns.regplot(x="engine-size", y="price", data=df)
plt.ylim(0,)
```

Engine-size

Correlation Coefficient: 0.8723351674455185 P-value: 9.265491622198389e-64

```
[11]: (0.0, 53414.609178181076)
```




```
[12]: print('Horsepower')
pearson_coef, p_value = stats.pearsonr(df['horsepower'], df['price'])
print("Correlation Coefficient:", pearson_coef, " P-value:", p_value)
print(" ")
#Since the p-value is < 0.001, the correlation between horsepower and price is
↳statistically
#significant, and the linear relationship is quite strong (~0.809, close to 1).

sns.regplot(x="horsepower", y="price", data=df)
plt.ylim(0,)
```

Horsepower

Correlation Coefficient: 0.809574567003656 P-value: 6.369057428259557e-48

```
[12]: (0.0, 47540.810858772275)
```




```
[13]: print('City-mpg')
pearson_coef, p_value = stats.pearsonr(df['city-mpg'], df['price'])
print("Correlation Coefficient:", pearson_coef, " P-value:", p_value)
print(" ")
#Since the p-value is < 0.001, the correlation between city-mpg and price is
↪statistically
#significant, and the coefficient of about -0.687 shows that the relationship
↪is negative and
#moderately strong.

sns.regplot(x="city-mpg", y="price", data=df)
plt.ylim(0,)
```

City-mpg

Correlation Coefficient: -0.6865710067844677 P-value: 2.321132065567674e-29

```
[13]: (0.0, 48246.761410914216)
```




```
[14]: print('Highway-mpg')
pearson_coef, p_value = stats.pearsonr(df['highway-mpg'], df['price'])
print( "Correlation Coefficient:", pearson_coef, " P-value:", p_value)
print(" ")
#Since the p-value is < 0.001, the correlation between highway-mpg and price is
↳statistically
#significant, and the coefficient of about -0.705 shows that the relationship
↳is negative and
#moderately strong.

sns.regplot(x="highway-mpg", y="price", data=df)
plt.ylim(0,)
```

Highway-mpg

Correlation Coefficient: -0.7046922650589529 P-value: 1.7495471144477352e-31

```
[14]: (0.0, 48160.81853259754)
```




```
[15]: print("Wheel-base")
pearson_coef, p_value = stats.pearsonr(df['wheel-base'], df['price'])
print("Correlation Coefficient:", pearson_coef, " P-value:", p_value)
print(" ")
#Since the p-value is < 0.001, the correlation between wheel-base and price is
↳ statistically
#significant, although the linear relationship isn't extremely strong (~0.585).

sns.regplot(x="wheel-base", y="price", data=df)
plt.ylim(0,)
```

Wheel-base

Correlation Coefficient: 0.5846418222655081 P-value: 8.076488270732989e-20

```
[15]: (0.0, 47614.140319248014)
```




```
[16]: print('Bore')
pearson_coef, p_value = stats.pearsonr(df['bore'], df['price'])
print("Correlation Coefficient:", pearson_coef, " P-value:", p_value)
#Since the p-value is < 0.001, the correlation between bore and price is
↪statistically
#significant, but the linear relationship is only moderate (~0.521)

sns.regplot(x="bore", y="price", data=df)
plt.ylim(0,)
```

Bore

Correlation Coefficient: 0.5431553832626602 P-value: 8.049189483935489e-17

```
[16]: (0.0, 47804.113699511036)
```



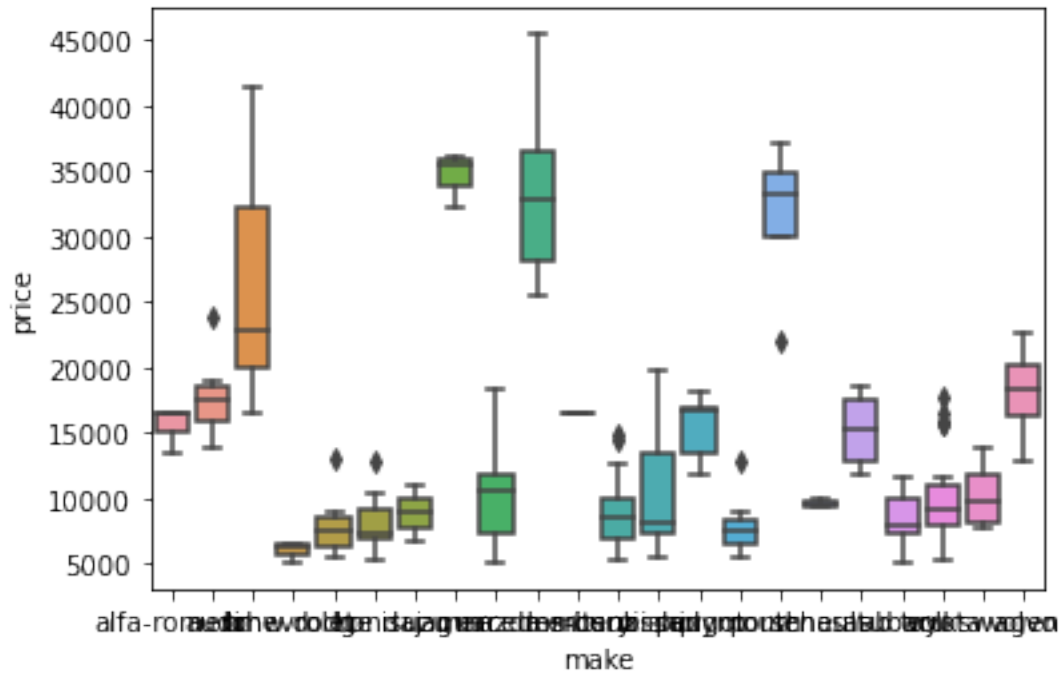
Categorical variables

```
[18]: #make, aspiration, num-of-doors, body-style, drive-wheels, engine-location,
      ↪ engine-type, num-of-cylinders, #fuel-system

      #Exclude variables that are not appropriate for predictors
```


```
[19]: #Make
      sns.boxplot(x="make", y="price", data=df)
      #exclude due to overlapping
```

```
[19]: <AxesSubplot:xlabel='make', ylabel='price'>
```


```
[20]: #Num-of-doors
sns.boxplot(x="num-of-doors", y="price", data=df)
#exclude due to overlapping
```

```
[20]: <AxesSubplot:xlabel='num-of-doors', ylabel='price'>
```




```
[21]: #Body-style
sns.boxplot(x="body-style", y="price", data=df)
#exclude due to overlapping
```

```
[21]: <AxesSubplot:xlabel='body-style', ylabel='price'>
```



```
[22]: #Drive-wheels
sns.boxplot(x="drive-wheels", y="price", data=df)
```

```
[22]: <AxesSubplot:xlabel='drive-wheels', ylabel='price'>
```



```
[23]: #Count per drive-wheels
drive_wheels_counts = df['drive-wheels'].value_counts().to_frame()
drive_wheels_counts.rename(columns={'drive-wheels': 'value_counts'},
                             inplace=True)
drive_wheels_counts.index.name = 'drive-wheels'
drive_wheels_counts
```

```
[23]:          value_counts
drive-wheels
fwd              118
rwd              75
4wd               8
```

```
[24]: #Average price for drive-wheels
df_group_one = df[['drive-wheels', 'body-style', 'price']]
df_group_one = df_group_one.groupby(['drive-wheels'], as_index=False).mean()
df_group_one.rename(columns={'price': 'avg price'}, inplace=True)
df_group_one
```

```
[24]:  drive-wheels    avg price
0         4wd  10241.000000
1         fwd   9244.779661
2         rwd  19757.613333
```

```
[25]: #Average price for drive-wheels and body-style
df_gptest = df[['drive-wheels','body-style','price']]
grouped_test1 = df_gptest.groupby(['drive-wheels','body-style'],
    ↪as_index=False).mean()
grouped_pivot = grouped_test1.pivot(index='drive-wheels', columns='body-style')
grouped_pivot = grouped_pivot.fillna(0)
grouped_pivot
```

```
[25]:
```

| | price | | | |
|--------------|-------------|--------------|--------------|--------------|
| body-style | convertible | hardtop | hatchback | sedan |
| drive-wheels | | | | |
| 4wd | 0.0 | 0.000000 | 7603.000000 | 12647.333333 |
| fwd | 11595.0 | 8249.000000 | 8396.387755 | 9811.800000 |
| rwd | 23949.6 | 24202.714286 | 14337.777778 | 21711.833333 |

| body-style | wagon |
|--------------|--------------|
| drive-wheels | |
| 4wd | 9095.750000 |
| fwd | 9997.333333 |
| rwd | 16994.222222 |

```
[26]: #Average price for body-style, heatmap to better understand distribution
fig, ax = plt.subplots()
im = ax.pcolor(grouped_pivot, cmap='RdBu')


#Label names
row_labels = grouped_pivot.columns.levels[1]
col_labels = grouped_pivot.index

#Move ticks and labels to the center
ax.set_xticks(np.arange(grouped_pivot.shape[1]) + 0.5, minor=False)
ax.set_yticks(np.arange(grouped_pivot.shape[0]) + 0.5, minor=False)

#Insert labels
ax.set_xticklabels(row_labels, minor=False)
ax.set_yticklabels(col_labels, minor=False)

#Rotate label if too long
plt.xticks(rotation=90)

fig.colorbar(im)
plt.show()
```



```
[28]: #ANOVA for drive-wheels
grouped_test2=df_gptest[['drive-wheels', 'price']].groupby(['drive-wheels'])

f_val, p_val = stats.f_oneway(grouped_test2.get_group('fwd')['price'],
    ↳grouped_test2.get_group('rwd')['price'], grouped_test2.
    ↳get_group('4wd')['price'])
print( "ANOVA results: F =", f_val, "| P =", p_val)

#A large F-test score shows a strong correlation and a P-value of almost 0
↳implies almost certain statistical significance.
```

ANOVA results: F = 67.95406500780399 | P = 3.3945443577151245e-23

```
[29]: #Check correlations of particular groups:

#ANOVA for drive-wheels / fwd and rwd
f_val, p_val = stats.f_oneway(grouped_test2.get_group('fwd')['price'],
    ↳grouped_test2.get_group('rwd')['price'])

print( "ANOVA results: F =", f_val, "| P =", p_val )
```

ANOVA results: F = 130.5533160959111 | P = 2.2355306355677845e-23

```
[30]: #ANOVA for drive-wheels / 4wd and rwd
f_val, p_val = stats.f_oneway(grouped_test2.get_group('4wd')['price'],
    ↪grouped_test2.get_group('rwd')['price'])

print( "ANOVA results: F =", f_val, "| P =", p_val)
```

ANOVA results: F = 8.580681368924756 | P = 0.004411492211225333


```
[31]: #ANOVA for drive-wheels / 4wd and fwd
f_val, p_val = stats.f_oneway(grouped_test2.get_group('4wd')['price'],
    ↪grouped_test2.get_group('fwd')['price'])

print("ANOVA results: F =", f_val, "| P =", p_val)
```

ANOVA results: F = 0.665465750252303 | P = 0.41620116697845666

```
[32]: #Engine-location
sns.boxplot(x="engine-location", y="price", data=df)
```

[32]: <AxesSubplot:xlabel='engine-location', ylabel='price'>



```
[33]: #Count per engine-location
engine_loc_counts = df['engine-location'].value_counts().to_frame()
engine_loc_counts.rename(columns={'engine-location': 'value_counts'},
    ↪inplace=True)
```


```
engine_loc_counts.index.name = 'engine-location'
engine_loc_counts
```

```
#exclude because there are only 3 rears
```

```
[33]: value_counts
engine-location
front          198
rear            3
```

```
[34]: #Num-of-cylinders
sns.boxplot(x="num-of-cylinders", y="price", data=df)
```

```
[34]: <AxesSubplot:xlabel='num-of-cylinders', ylabel='price'>
```




```
[35]: #Count per num-of-cylinders
engine_loc_counts = df['num-of-cylinders'].value_counts().to_frame()
engine_loc_counts.rename(columns={'num-of-cylinders': 'value_counts'},
                           inplace=True)
engine_loc_counts.index.name = 'num-of-cylinders'
engine_loc_counts

#exclude because most models fall into four cylinders category
```

```
[35]: value_counts
num-of-cylinders
four      157
six       24
five      10
two        4
eight      4
three      1
twelve     1
```

```
[36]: #Fuel-system
sns.boxplot(x="fuel-system", y="price", data=df)
#exclude due to overlapping
```

```
[36]: <AxesSubplot: xlabel='fuel-system', ylabel='price'>
```



Conclusion

```
[38]: #The following should be considered as predictors:

#----Numerical----

# Length
# Width
# Curb-weight
```



```
#   Engine-size
#   Horsepower
#   City-mpg
#   Highway-mpg
#   Wheel-base
#   Bore

#----Categorical----

#   Drive-wheels
```